

QSentinel

Quantum-Inspired Cyber Threat Detection for Digital Signature Security

Non-ML Statistical Threat Detection for
Teleportation-Based Quantum Digital Signatures

Problem Statement ID	26141
Team ID	135012
Team Name	OK Computer
Theme	Blockchain and Cybersecurity
Competition	Smart India Hackathon 2026

Team Members

Tanisha Deka
Dwibon Bhargab Deka
Barnil Mahanta
Nayan Nirban Dewri
Arpan Goswami
Swornima Dey

Abstract

Quantum Digital Signature (QDS) systems require mechanisms for detecting tampering with quantum communication channels in addition to conventional protocol-level threats such as replay, impersonation, and unauthorized verification. This work presents **QSentinel**, a simulation-based, non-machine-learning threat-detection framework for a teleportation-based QDS setting. The proposed approach embeds a small, secret, session-bound set of Pauli-eigenstate sentinel qubits within the transmitted sequence. Measurements of these sentinels produce an axis-resolved disturbance fingerprint (e_X, e_Y, e_Z) , allowing the system to distinguish directional channel disturbances that may be obscured by a single aggregate error rate.

QSentinel first calibrates the expected error level of the honest channel and then applies exact one-sided binomial hypothesis tests independently to the three Pauli axes. The resulting statistical evidence is combined with protocol-level checks for session freshness, replay, verifier authorization, signature consistency, and teleportation consistency. The detector produces three outcomes: **ACCEPT**, **INCONCLUSIVE**, and **REJECT**, with multiple-axis evidence required for rejection.

The framework was evaluated through simulations of X-, Y-, and Z-Pauli attacks, intercept-resend, Z-measure-resend, partial attacks, and channel noise drift. In the recorded 100-trial experiments, all five principal full-strength attack classes achieved 100% observed detection. A separate 10,000-trial honest-channel experiment produced 9,728 accepted, 271 inconclusive, and 1 rejected trial. These results demonstrate the feasibility of using calibrated, axis-resolved statistical evidence for interpretable quantum-channel threat detection under the tested assumptions.

QSentinel is implemented as an end-to-end software prototype comprising the quantum simulation and detection core, protocol logic, automated tests, FastAPI backend, and web-based dashboard. The work is intended as a practical detection layer and experimental framework rather than a complete production QDS protocol or a formal unconditional or composable security proof.

Contents

Abstract	1
1 Problem and Proposed Solution	4
1.1 Problem	4
1.2 Solution at a Glance	4
1.3 Why the Approach is Non-ML	4
2 Quantum and Protocol Foundations	5
2.1 Pauli Eigenstates and the Fingerprint	5
2.2 Teleportation Model	5
2.3 Protocol Security Controls	5
3 QSentinel Detection Method	7
3.1 Secret Session-Bound Sentinels	7
3.2 Measurement and Error Rates	7
3.3 Statistical Decision Rule	7
3.4 Meaning of INCONCLUSIVE	7
4 Attack Models	8
5 Implementation and Public Prototype	9
5.1 Software Architecture	9
5.2 Core Repository Structure	9
5.3 Backend and Dashboard	9
5.4 Deployment	9
6 Experimental Evaluation	11
6.1 Evaluation Configuration	11
6.2 Representative Fingerprints	11
6.3 Full-Strength Attack Detection	11
6.4 Honest-Channel Validation	11
6.5 Wrong-Key Verification	11
6.6 Partial-Attack Power	12
6.7 Noise Drift	12
6.8 Fingerprint Visualization	13
7 Security Interpretation and Limits	14
7.1 Why Axis-Resolved Evidence Helps	14
7.2 Identifiability Limit	14

7.3	Security Assumptions	14
7.4	What QSentinel Does Not Claim	14
7.5	Holevo Information: Correct Scope	14
8	Testing, Reproducibility and Limitations	15
8.1	Automated Testing	15
8.2	Reproduce the Prototype	15
8.3	Current Limitations	15
9	Current Contribution and Future Work	16
9.1	Contribution Positioning	16
9.2	Future Work	16
10	Conclusion	17
	References	18

1 Problem and Proposed Solution

1.1 Problem

Quantum Digital Signatures require more than ordinary message-integrity checks. In the target teleportation-based setting, the system must detect or flag quantum-channel manipulation while also handling classical protocol threats such as forgery, impersonation, replay and unauthorized verification.

A single error rate is useful for detecting disturbance but loses directional information. Different attacks can produce different patterns across the three Pauli axes. The problem addressed by **QSentinel** is therefore:

How can a software system obtain interpretable quantum evidence of channel tampering, compare it against an honest noise floor, and combine that evidence with protocol-level security checks without using machine learning?

1.2 Solution at a Glance

The proposed workflow is:

1. generate a signature/session and derive a secret, session-bound sentinel schedule;
2. transmit ordinary payload states together with hidden X/Y/Z Pauli-eigenstate sentinels;
3. model teleportation, channel noise and attack behaviour;
4. measure sentinels and estimate (e_X, e_Y, e_Z) ;
5. compare each axis against its calibrated honest baseline using an exact one-sided binomial test;
6. combine the quantum result with replay, authorization, freshness and signature checks;
7. expose the result through a web dashboard and log security events.




Figure 1: QSentinel end-to-end detection workflow.

1.3 Why the Approach is Non-ML

The detector does not train a classifier. Its evidence comes directly from measured sentinel outcomes, a calibrated null model, and an explicit statistical decision rule. This makes the reason for a decision inspectable: the system can show which axis exceeded the expected honest-channel error level and why the result became inconclusive or reject.

2 Quantum and Protocol Foundations

2.1 Pauli Eigenstates and the Fingerprint

The Pauli operators are

$$X = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad Y = \begin{pmatrix} 0 & -i \\ i & 0 \end{pmatrix}, \quad Z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}.$$

Their eigenstates define the X, Y and Z measurement axes. A Pauli operation leaves its own eigenstates unchanged up to phase while disturbing eigenstates of the other axes. Consequently, strong Pauli attacks produce characteristic fingerprints:

Table 1: Idealized error fingerprints without channel noise.

Condition	e_X	e_Y	e_Z
None	0	0	0
X-Pauli	0	1	1
Y-Pauli	1	0	1
Z-Pauli	1	1	0
Intercept-resend	$\approx 1/3$	$\approx 1/3$	$\approx 1/3$
Z-measure-resend	$\approx 1/2$	$\approx 1/2$	≈ 0

These are ideal fingerprints, not guaranteed experimental values. Honest channel noise shifts the measured rates upward, which is why QSentinel calibrates a baseline before detection.

2.2 Teleportation Model

The simulator follows the standard teleportation structure: an input state is combined with an entangled Bell pair, a Bell measurement yields two classical bits, and the receiver applies the corresponding Pauli correction. In compact form, the pre-correction receiver state is represented by

$$X^{m_2} Z^{m_1} |\psi\rangle, \quad (m_1, m_2) \in \{0, 1\}^2.$$



Figure 2: Simplified teleportation structure represented in the simulator.

2.3 Protocol Security Controls

The prototype separates quantum evidence from protocol evidence:

- **Replay:** a previously used session/signature is rejected rather than treated as a new transmission.
- **Verifier authorization:** the verifier identity must match the registered verifier for the session.
- **Freshness:** the sentinel schedule is derived from secret key material and the session identifier using HMAC-SHA256.
- **Wrong-key verification:** verification with incorrect key material produces strong sentinel disturbance in the tested setting.

- **Teleportation consistency:** Bell outcomes and Pauli corrections are part of the simulated workflow.

3 QSentinel Detection Method

3.1 Secret Session-Bound Sentinels

For each session, QSentinel derives the sentinel positions from secret key material and a session identifier. The derivation uses HMAC-SHA256 with the context string `QSentinel-Sentinel-Schedule-v2:`. Thus, the schedule is reproducible for the authorized session but changes when the session changes. A recorded freshness check gave:

$$A = B : \text{False}, \quad A = C : \text{True},$$

where A and B use different session identifiers and C repeats A's session.

3.2 Measurement and Error Rates

Let n_a be the number of sentinels measured on axis $a \in \{X, Y, Z\}$ and k_a the number of errors. The empirical error rate is

$$\hat{e}_a = \frac{k_a}{n_a}.$$

The resulting vector $(\hat{e}_X, \hat{e}_Y, \hat{e}_Z)$ is the quantum fingerprint used by the detector.

For an honest depolarizing channel with parameter p , the simulated per-axis error probability is

$$\theta = \frac{2p}{3}.$$

The system estimates the honest baseline through calibration sessions rather than assuming the theoretical value is exact.

3.3 Statistical Decision Rule

For each axis, the detector tests

$$H_0 : e_a \leq b_a \quad \text{vs.} \quad H_1 : e_a > b_a,$$

where b_a is the calibrated honest baseline. An exact one-sided binomial p-value is compared with significance level α .

The final detector is intentionally conservative:

- 0 rejected axes $\rightarrow$ **ACCEPT**;
- 1 rejected axis $\rightarrow$ **INCONCLUSIVE**;
- 2 or more rejected axes $\rightarrow$ **REJECT**.

This design avoids treating a single noisy axis as sufficient evidence of an attack.

3.4 Meaning of INCONCLUSIVE

INCONCLUSIVE is a deliberate third outcome, not a failure of the detector. It means the observed evidence is abnormal enough to require caution but does not satisfy the multi-axis rejection rule. This is particularly relevant for weak attacks and channel drift.

4 Attack Models

The simulator implements both directional and basis-independent disturbance models.

Table 2: Attack models used in evaluation.

Attack	Modelled behaviour
X/Y/Z-Pauli	Applies the corresponding Pauli operation to targeted states.
Intercept-resend	Intercepts and resends according to a measurement/resend model, producing approximately balanced disturbance across axes.
Z-measure-resend	Measures in the Z basis and resends the resulting state; X/Y sentinels are strongly disturbed while Z is comparatively preserved.
Partial attack	Applies an attack with configurable strength to study statistical detection power.
Depolarizing noise	Represents honest-channel noise and controlled noise drift.

The attack models are deliberately explicit and interpretable. They are not intended to represent every possible adaptive quantum adversary.

5 Implementation and Public Prototype

5.1 Software Architecture

The implementation is organized into separable modules for quantum states, attacks, sentinel scheduling and measurement, noise, detection, teleportation, protocol logic and the QDS workflow. Experiments are kept separately so that calibration, robustness and statistical validation can be reproduced independently.




Figure 3: Software architecture of the prototype.

5.2 Core Repository Structure

```

QSentinel/
  src/ states, attacks, sentinels, detector,
        teleportation, protocol, QDS pipeline
  experiments/ calibration, attack, robustness and evaluation scripts
  backend/ FastAPI application and deployment requirements
  QSentinel_Frontend_Animated/
        React dashboard
  tests/ automated tests
  
```

5.3 Backend and Dashboard

The backend exposes signature creation and verification, replay simulation, attack simulation, forgery/impersonation scenarios and security-event logging. The public dashboard provides signature creation, clean verification, replay testing and a Threat Lab for attack scenarios.

Table 3: Main public API operations.

Route	Purpose
/signature/create	Create a simulated signature/session
/signature/verify	Verify a signature
/signature/replay	Test replay handling
/simulate	Run an attack/noise scenario
/simulate/forgery	Wrong-key forgery simulation
/simulate/impersonation	Impersonation scenario
/events	Security-event logging

5.4 Deployment

The public prototype uses a GitHub Pages React frontend and a Render-hosted FastAPI backend. The backend is a free-tier deployment and may require a short wake-up period after inactivity.




Figure 4: Public deployment architecture.

Public resources:

<https://dwibon.github.io/QSentinel/>

<https://github.com/Dwibon/QSentinel>

Demo note. Because the backend uses Render’s free tier, the first request after inactivity can take approximately 50 seconds to 1 minute. If the dashboard appears inactive, wait for the backend to wake and then retry.

6 Experimental Evaluation

6.1 Evaluation Configuration

Unless otherwise stated, the main quantitative experiments used $n = 5000$ qubits, a sentinel fraction of 20%, honest-channel noise $p = 0.03$, baseline 0.02 per axis, and $\alpha = 0.01$. Detection-rate experiments used 100 trials per condition.

6.2 Representative Fingerprints

A representative integrated run produced:

Table 4: Representative end-to-end fingerprints.

Condition	e_X	e_Y	e_Z	Decision
None	.020	.028	.021	ACCEPT
X-Pauli	.009	.981	.988	REJECT
Y-Pauli	.976	.021	.962	REJECT
Z-Pauli	.979	.968	.006	REJECT
Intercept-resend	.344	.301	.366	REJECT
Z-measure-resend	.522	.500	.022	REJECT

6.3 Full-Strength Attack Detection

For X-Pauli, Y-Pauli, Z-Pauli, intercept-resend and Z-measure-resend, each condition achieved 100/100 observed detection in the recorded 100-trial experiments. For a 100/100 result, the two-sided 95% Clopper–Pearson lower bound is approximately 96.38%; the result is therefore empirical rather than a universal guarantee.

6.4 Honest-Channel Validation

A separate 10,000-trial honest validation gave:

Table 5: 10,000-trial honest-channel validation.

Outcome	Count	Observed rate
ACCEPT	9,728	97.28%
INCONCLUSIVE	271	2.71%
REJECT	1	0.01%

The observed rejection rate was $1/10,000 = 0.01\%$. This is an observed simulation rate, not a theoretical false-positive guarantee. The 95% upper bound for the true rejection probability under this tested condition is approximately 5.57×10^{-4} .

6.5 Wrong-Key Verification

In 100 wrong-key trials, the mean fingerprint was

$$(0.4985, 0.5069, 0.5020),$$

with 100 REJECT, 0 INCONCLUSIVE and 0 ACCEPT. This experiment demonstrates wrong-key verification behaviour; it should not be interpreted as a complete proof of signer impersonation resistance.

6.6 Partial-Attack Power

Table 6: Observed detection power for partial attacks; 100 trials per condition.

Strength	X	Y	Z	Intercept	Z-measure
0.01	.020	.000	.000	.000	.000
0.02	.190	.140	.190	.000	.000
0.05	.910	.900	.930	.150	.430
0.10	1.000	1.000	1.000	.930	.950
0.20+	1.000	1.000	1.000	1.000	1.000

The results illustrate the statistical nature of the system: weak attacks can remain below the rejection threshold, while stronger attacks become increasingly detectable.

6.7 Noise Drift

With a calibrated baseline of 0.02, 100 honest trials were run at each noise level:

Table 7: Observed decisions under honest-channel noise drift.

p	ACCEPT	INCONCLUSIVE	REJECT
0.020	100		0
0.022	100		0
0.025	99		1
0.030	94		6
0.035	95		5
0.040	79		20

The main observed effect of moderate drift was movement toward **INCONCLUSIVE**; larger drift eventually produced rejection. This motivates periodic recalibration.

6.8 Fingerprint Visualization




Figure 5: Representative axis-resolved fingerprints used by the detector.

7 Security Interpretation and Limits

7.1 Why Axis-Resolved Evidence Helps

The vector (e_X, e_Y, e_Z) preserves directional information that an aggregate error rate discards. A strong X-Pauli attack, for example, leaves the X eigenstates comparatively stable while strongly disturbing Y and Z sentinels. This makes the evidence interpretable rather than merely anomalous.

7.2 Identifiability Limit

For a general channel represented by a Bloch-vector map T , the error observed by a Pauli-axis eigenstate satisfies

$$e_a = \frac{1 - T_{aa}}{2}.$$

Therefore, two different channels with the same diagonal elements can produce the same three sentinel error rates. The fingerprint is consequently a useful detector feature, not a complete characterization of an arbitrary quantum channel.

7.3 Security Assumptions

The prototype assumes that the secret key material and session binding used for sentinel scheduling are protected, the authorized verifier can identify the expected sentinel bases, and the statistical null model reasonably represents honest operation. It also assumes the tested attack/channel models are relevant to the scenario being evaluated.

7.4 What QSentinel Does Not Claim

QSentinel does not provide a universal detector for all quantum attacks, a composable security proof, a complete production QDS/QKD protocol, or a hardware-level security guarantee. Replay and verifier authorization are protocol controls rather than direct quantum-channel measurements. Experimental detection rates are observed simulation results.

7.5 Holevo Information: Correct Scope

Holevo information can be used to quantify an upper bound on accessible classical information about a quantum ensemble under stated assumptions. It is not itself a detector and does not by itself provide a forging probability. Any stronger claim about guessing probability requires an explicit min-entropy or equivalent security analysis.

8 Testing, Reproducibility and Limitations

8.1 Automated Testing

The recorded test suite contains 78 passing tests. It covers state preparation and measurement, attacks, noise, sentinel scheduling, detection decisions, teleportation/protocol logic and API-facing behaviour. The freshness test confirms that different session identifiers produce different schedules while the same session reproduces the schedule.

8.2 Reproduce the Prototype

```
git clone https://github.com/Dwibon/QSentinel.git
cd QSentinel
python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt
python -m pytest tests/

# backend
uvicorn backend.app:app --reload

# frontend, in another terminal
cd QSentinel_Frontend_Animated
npm install
npm run dev
```

Local backend: <http://127.0.0.1:8000>.

Local frontend: <http://localhost:5173>.

8.3 Current Limitations

- The quantum processor is simulated; results are not hardware measurements.
- The signature construction is a prototype abstraction rather than a complete production QDS construction.
- The integer key represents secret material; a complete BB84/QKD protocol is not implemented.
- Detection depends on calibration and can become inconclusive or inaccurate under substantial noise drift.
- The evaluated attack set is finite; adaptive or correlated attacks may behave differently.
- The current detector does not establish a formal security theorem.

9 Current Contribution and Future Work

9.1 Contribution Positioning

QSentinel builds on established concepts from quantum teleportation, quantum digital signatures, quantum authentication/test-state methods, Pauli operators and statistical inference. The project does not claim these individual components as novel.

The integrated contribution is a system-level methodology that combines:

1. secret session-bound sentinel scheduling;
2. three-axis quantum disturbance fingerprints;
3. honest-channel calibration;
4. exact per-axis statistical testing with a conservative multi-axis rule;
5. explicit simulation of several attack classes and partial attack strength;
6. protocol-level replay and authorization checks;
7. an end-to-end API and dashboard suitable for reproducible demonstration.

9.2 Future Work

- Validation on actual quantum hardware and broader quantum-channel models.
- Evaluation against more sophisticated correlated, adaptive, and coherent attacks.
- Formal security analysis of detection guarantees and error thresholds.
- Adaptive calibration to handle changing noise levels and channel drift.
- Larger-scale statistical evaluation across more qubits and repeated trials.
- Integration with a complete QDS protocol implementation.
- Hardware-backed verifier authentication.
- Extension to additional quantum observables beyond the three Pauli axes.

10 Conclusion

QSentinel addresses the practical detection problem in the SIH setting by turning hidden Pauli-eigenstate test transmissions into an interpretable three-axis statistical fingerprint. Calibration makes the detector sensitive to the honest noise floor, while the **INCONCLUSIVE** state avoids forcing every borderline observation into accept or reject. Protocol checks cover threats that quantum-channel statistics alone cannot establish, such as replay and verifier authorization.

The current prototype demonstrates the complete path from signature/session generation through teleportation and attack simulation to statistical detection, API exposure, dashboard visualization and logging. Its experimental results support the feasibility of the proposed methodology under the tested assumptions while also defining clear limits on what can be claimed.

References

- [1] D. Gottesman and I. L. Chuang, “Quantum Digital Signatures,” arXiv:quant-ph/0105032, 2001.
- [2] P. Wallden, V. Dunjko, A. Kent, and E. Andersson, “Quantum digital signatures with quantum-key-distribution components,” *Physical Review A*, vol. 91, 042304, 2015.
- [3] C. H. Bennett, G. Brassard, C. Crépeau, R. Jozsa, A. Peres, and W. K. Wootters, “Teleporting an unknown quantum state via dual classical and Einstein-Podolsky-Rosen channels,” *Physical Review Letters*, vol. 70, pp. 1895–1899, 1993.
- [4] H. Barnum, C. Crépeau, D. Gottesman, A. Smith, and A. Tapp, “Authentication of quantum messages,” in *Proceedings of the 43rd IEEE Symposium on Foundations of Computer Science*, 2002.
- [5] M. H. Saeed, H. Sattar, M. H. Durad, and Z. Haider, “Implementation of QKD BB84 Protocol in Qiskit,” in *Proceedings of the 19th International Bhurban Conference on Applied Sciences and Technology (IBCAST)*, 2022, pp. 689–695, doi: 10.1109/IBCAST54850.2022.9990073.
- [6] A. Tarable, R. P. Paganelli, E. Storelli, A. Gatto, and M. Ferrari, “Generalized Quantum-assisted Digital Signature,” arXiv:2406.19978, 2024.
- [7] D. Lu, Z. Li, J. Yu, and Z. Han, “A Verifiable Arbitrated Quantum Signature Scheme Based on Controlled Quantum Teleportation,” *Entropy*, vol. 24, no. 1, 111, 2022, doi: 10.3390/e24010111.
- [8] X. Tan and T.-Y. Ye, “Semiquantum proxy blind signature based on quantum teleportation,” 2022.
- [9] J. W. Choi, K.-Y. Chang, and D. Hong, “Security problem on arbitrated quantum signature schemes,” *Physical Review A*, vol. 84, 062330, 2011.
- [10] P.-A. Jacqmin and J. Liénardy, “Cryptanalysis of four arbitrated quantum signature schemes,” arXiv:2603.19985, 2026.
- [11] G. Xu, S. Xue, X. Zou, and Y. Zhang, “On a Grover-Based Quantum Signature Scheme and a Teleportation-Based Design,” *Entropy*, vol. 28, no. 9, 1030, 2026, doi: 10.3390/e28091030.
- [12] Qniverse, “Quantum Digital Signature (QDS) Algorithm,” online tutorial, <https://qniverse.in/docs/quantum-digital-signatureqds-algorithm/>.
- [13] E. B. Wilson, “Probable inference, the law of succession, and statistical inference,” *Journal of the American Statistical Association*, vol. 22, no. 158, pp. 209–212, 1927.
- [14] W. Hoeffding, “Probability inequalities for sums of bounded random variables,” *Journal of the American Statistical Association*, vol. 58, no. 301, pp. 13–30, 1963.